

WP4 update - Analytical model demonstration

22nd March 2024



This project has received funding from the European Union's Horizon Research and Innovation Actions under Grant Agreement № 101093202.

Naïve sequential BFS algorithm

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
    num_nodes)
2     for (int i = 0; i < num_nodes; ++i)
3         level[i] = -1;
4         parent[i] = -1;
5
6     std::queue<int> queue;
7
8     level[src] = 0;
9     queue.push(src);
10
11    while (!queue.empty())
12        int u = queue.front();
13        queue.pop();
14
15        for (int v = 0; v < num_nodes; ++v)
16            if (G[u][v] > 0 && level[v] == -1)
17                parent[v] = u;
18                level[v] = level[u] + 1;
19                queue.push(v);
```

Assumptions

- Graph only has a single connected component
- Write back cache
 - Memory write takes the same time as memory read

Constructing symbolical model

$$T_{BFS} = nT_{init}$$

$$T_{init} = 2T_{mem_write}$$

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
  num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6   std::queue<int> queue;
7
8   level[src] = 0;
9   queue.push(src);
10
11  while (!queue.empty())
12    int u = queue.front();
13    queue.pop();
14
15    for (int v = 0; v < num_nodes; ++v)
16      if (G[u][v] > 0 && level[v] == -1)
17        parent[v] = u;
18        level[v] = level[u] + 1;
19        queue.push(v);
```

Constructing symbolical model

$$T_{BFS} = nT_{init}$$

$$T_{init} = 2T_{mem_write}$$

We skip this!

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
  num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6     std::queue<int> queue;
7
8     level[src] = 0;
9     queue.push(src);
10
11    while (!queue.empty())
12      int u = queue.front();
13      queue.pop();
14
15      for (int v = 0; v < num_nodes; ++v)
16        if (G[u][v] > 0 && level[v] == -1)
17          parent[v] = u;
18          level[v] = level[u] + 1;
19          queue.push(v);
```

Constructing symbolical model

$$T_{BFS} = nT_{init} + nT_{while_loop}$$

$$T_{init} = 2T_{mem_write}$$

$$T_{while_loop} = T_{q_front} + T_{q_pop} + n(T_{add} + T_{G_mem_read} + T_{mem_read})$$

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
  num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6   std::queue<int> queue;
7
8   level[src] = 0;
9   queue.push(src);
10
11  while (!queue.empty())
12    int u = queue.front();
13    queue.pop();
14
15    for (int v = 0; v < num_nodes; ++v)
16      if (G[u][v] > 0 && level[v] == -1)
17        parent[v] = u;
18        level[v] = level[u] + 1;
19        queue.push(v);
```

Constructing symbolical model

$$T_{BFS} = nT_{init} + nT_{while_loop} + nT_{visit_node}$$

$$T_{init} = 2T_{mem_write}$$

$$T_{while_loop} = T_{q_front} + T_{q_pop} + n(T_{add} + T_{G_mem_read} + T_{mem_read})$$

$$T_{visit_node} = 2T_{DRAM} + T_{L1} + T_{add} + T_{q_push}$$

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
   num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6   std::queue<int> queue;
7
8   level[src] = 0;
9   queue.push(src);
10
11  while (!queue.empty())
12    int u = queue.front();
13    queue.pop();
14
15    for (int v = 0; v < num_nodes; ++v)
16      if (G[u][v] > 0 && level[v] == -1)
17        parent[v] = u;
18        level[v] = level[u] + 1;
19        queue.push(v);
```

Constructing symbolical model

$$T_{BFS} = nT_{init} + nT_{while_loop} + nT_{visit_node}$$

$$T_{init} = 2T_{mem_write}$$

$$T_{while_loop} = T_{q_front} + T_{q_pop} + n(T_{add} + T_{G_mem_read} + T_{mem_read})$$

$$T_{visit_node} = 2T_{DRAM} + T_{L1} + T_{add} + T_{q_push}$$

$$T_{mem_write} = T_{mem_read} = (1 - MR_{L1})T_{L1} + MR_{L1}(1 - MR_{L2})T_{L2} + MR_{L1}MR_{L2}(1 - MR_{L3})T_{L3} + MR_{L1}MR_{L2}MR_{L3}T_{DRAM}$$

$$MR_{L\{1,2,3\}} = \frac{1}{cacheLineSize_{L\{x\}}/sizeof(int)}$$

$$T_{G_mem_read} = (1 - G_MR_{L1})T_{L1} + G_MR_{L1}(1 - G_MR_{L2})T_{L2} + G_MR_{L1}G_MR_{L2}(1 - G_MR_{L3})T_{L3} + G_MR_{L1}G_MR_{L2}G_MR_{L3}T_{DRAM}$$

$$G_MR_{L\{1,2,3\}} = \frac{1}{cacheLineSize_{L\{x\}}/sizeof(float)}$$

```

1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
  num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6   std::queue<int> queue;
7
8   level[src] = 0;
9   queue.push(src);
10
11  while (!queue.empty())
12    int u = queue.front();
13    queue.pop();
14
15    for (int v = 0; v < num_nodes; ++v)
16      if (G[u][v] > 0 && level[v] == -1)
17        parent[v] = u;
18        level[v] = level[u] + 1;
19        queue.push(v);

```


Calibrating the model

- Cache line sizes are 64 bytes for L1, L2, and L3
- Memory latency (measured by using LMBench `lat_mem_rd`):

Memory level	latency (ns)
L1	1.26
L2	4.42
L3	20.9
DRAM	62.5

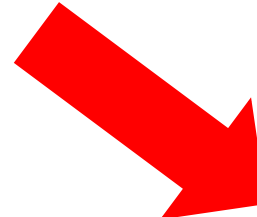
Table 1: Observed memory latency times

Calibrating the model

- Cache line sizes are 64 bytes for L1, L2, and L3
- Memory latency (measured by using LMBench `lat_mem_rd`):

Memory level	latency (ns)
L1	1.26
L2	4.42
L3	20.9
DRAM	62.5

Table 1: Observed memory latency times


$$MR_{L\{1,2,3\}} = G_MR_{L\{1,2,3\}} = \frac{1}{64/4} = \frac{1}{16}$$

$$\begin{aligned} T_{mem_write} &= T_{mem_read} = T_{G_mem_read} \\ &= \frac{15}{16} 1.26 + \frac{1}{16} \frac{15}{16} 4.42 + \frac{1}{16} \frac{1}{16} \frac{15}{16} 20.9 + \frac{1}{16} \frac{1}{16} \frac{1}{16} 62.5 \\ &= 1.521484375 \end{aligned}$$

Calibrating the model

- Operation latencies (measured by microbenchmarking):

Operation	latency (ns)
queue push	16.1
queue front	14.5
queue pop	11.2
integer add	0.326

Table 2: Observed operation latency times

Calibrating the model

- The final model:

$$T_{BFS} = nT_{init} + nT_{while_loop} + nT_{visit_node}$$

$$T_{init} = 2 \cdot 1.521484375$$

$$T_{while_loop} = 14.5 + 11.2 + n(0.326 + 1.521484375 + 1.521484375)$$

$$T_{visit_node} = 2 \cdot 62.5 + 1.26 + 0.326 + 16.1$$

$$\begin{aligned} T_{BFS} &= n(2 \cdot 1.521484375 + 14.5 + 11.2 + n(0.326 + 1.521484375 + 1.521484375) \\ &\quad + 2 \cdot 62.5 + 1.26 + 0.326 + 16.1) \\ &= 171.429n + 3.36897n^2 \end{aligned}$$

Calibrating the model

- The final model:

$$T_{BFS} = nT_{init} + nT_{while_loop} + nT_{visit_node}$$

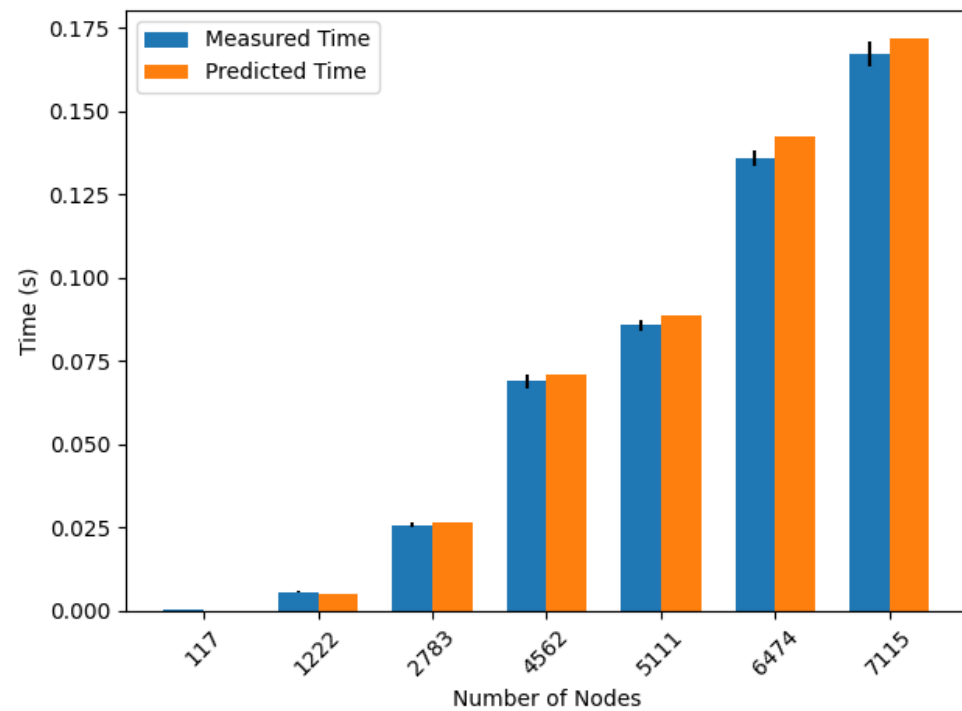
$$T_{init} = 2 \cdot 1.521484375$$

$$T_{while_loop} = 14.5 + 11.2 + n(0.326 + 1.521484375 + 1.521484375)$$

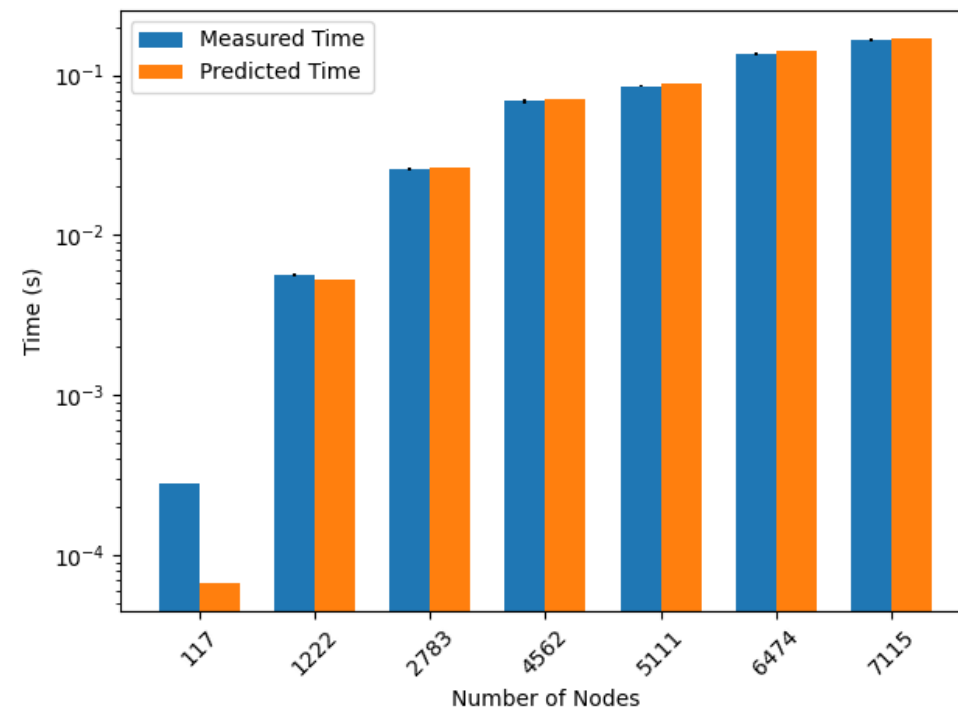
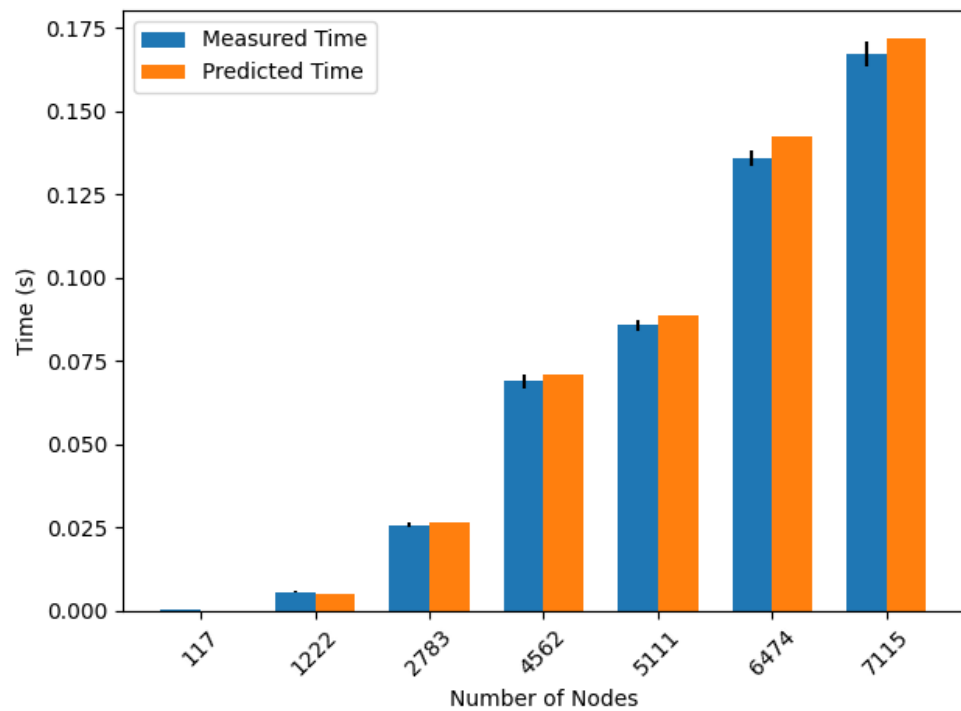
$$T_{visit_node} = 2 \cdot 62.5 + 1.26 + 0.326 + 16.1$$

$$\begin{aligned} T_{BFS} &= n(2 \cdot 1.521484375 + 14.5 + 11.2 + n(0.326 + 1.521484375 + 1.521484375)) \\ &\quad + 2 \cdot 62.5 + 1.26 + 0.326 + 16.1) \\ &= 171.429n + 3.36897n^2 \end{aligned}$$

Accuracy



Accuracy



Constructing symbolical model



$$T_{BFS} = nT_{init}$$

$$T_{init} = 2T_{mem_write}$$

We skip this!

```
1 void bfs(int *level, int *parent, std::vector<float> *G, int src, int
  num_nodes)
2   for (int i = 0; i < num_nodes; ++i)
3     level[i] = -1;
4     parent[i] = -1;
5
6     std::queue<int> queue;
7
8     level[src] = 0;
9     queue.push(src);
10
11    while (!queue.empty())
12      int u = queue.front();
13      queue.pop();
14
15      for (int v = 0; v < num_nodes; ++v)
16        if (G[u][v] > 0 && level[v] == -1)
17          parent[v] = u;
18          level[v] = level[u] + 1;
19          queue.push(v);
```


Accuracy



Accuracy

